

Lab: Importing Log Data into HDFS using Flume

About this Lab

Objective:	Import data from a log file into HDFS using Flume.
File locations:	/root/devph/labs/Lab3.3
Successful outcome:	The data in webtraffic.log on sandbox will be streamed into HDFS directory /user/root/flumedata/.
Before you begin:	Your HDP 2.3 cluster should be up and running within your VM.
Related lesson:	<i>Flume</i>

Lab Steps

- 1) Verify Flume is installed on your cluster.
 - a. If not already done, open a Terminal in your VM and type "ssh sandbox".
 - b. Type the following command from prompt:

```
# flume-ng
```

- c. The command should return the following usage instructions:

```
[root@node1 ~]# flume-ng
Error: Unknown or unspecified command ''

Usage: /usr/hdp/2.2.0.0-2041/flume/bin/flume-ng.distro <command> [options]...

commands:
  help          display this help text
  agent         run a Flume agent
  avro-client   run an avro Flume client
  password      create a password file for use in flume config
  version       show Flume version info

global options:
  --conf, -c <conf>    use configs in <conf> directory
  --classpath, -C <cp> append to the classpath
  --dryrun, -d         do not actually start Flume, just print the command
  --plugins-path <dirs> colon-separated list of plugins.d directories. See the
                        plugins.d section in the user guide for more details.
                        Default: $FLUME_HOME/plugins.d
  -Dproperty=value     sets a Java system property value
  -Xproperty=value     sets a Java -X option

agent options:
  --conf-file, -f <file> specify a config file (required)
```

- 2) View the contents of the webtraffic.log file

- a. Change directories to /root/devph/labs/Labs3.3.

```
# cd /root/devph/labs/Lab3.3
```

- b. View the contents of the webtraffic.log:

```
# less webtraffic.log
```

- 3) Press q to quit viewing the log.

- a. A partial agent configuration has been written for you, view the file using the less command.

```
# less logagent.conf
```

- b. From the VM desktop, start gedit.
- c. Click Open and navigate to /root/devph/labs/Lab3.3
- d. Modify this file replacing strings starting with "REPLACE-WITH-" with their appropriate value. HINT: Search for "HDFS Sink" at <https://flume.apache.org/FlumeUserGuide.html>.

```
agent.sources = weblog
agent.channels = memoryChannel
agent.sinks = mycluster

## Sources
#####
agent.sources.weblog.type = exec
agent.sources.weblog.command = tail -F REPLACE-WITH-PATH2-
webtraffic.log-FILE
agent.sources.weblog.batchSize = 1
agent.sources.weblog.channels = REPLACE-WITH-CHANNEL-NAME

## Channels
#####
agent.channels.memoryChannel.type = memory
agent.channels.memoryChannel.capacity = 100
agent.channels.memoryChannel.transactionCapacity = 100

## Sinks
#####
agent.sinks.mycluster.type = REPLACE-WITH-CLUSTER-TYPE
agent.sinks.mycluster.hdfs.path=/user/root/flumedata
agent.sinks.mycluster.channel = REPLACE-WITH-CHANNEL-NAME
```

- ```
flume-ng agent -n agent -f logagent.conf &
```

- ```
# hdfs dfs -ls flumedata/
```

- ```
hdfs dfs -cat flumedata/FlumeData.<#####>
```

```
[root@node1 flume]# hdfs dfs -cat flumedata/FlumeData.1445290807809
SEQ[1]prg.apache.hadoop.io.LongWritable*org.apache.hadoop.io.BytesWritableSL*!84668617,http://www.hortonworks.com!874679672,http://www.facebook.com!82587230,http://cnn.com!80125703,http://cnn.com!89467386,http://cnn.com!581842607,http://www.facebook.com!8493259972,http://www.hortonworks.com!8607729813,http://www.hortonworks.com!853802950,http://www.yahoo.com!819561879,http://www.yahoo.com[root@node1 flume]#
```

## Result

3